


Transaction

Outline



- ☐ Transaction Concept
- ☐ Transaction State
- ☐ Concurrent Executions
- ☐ Serializability
- ☐ Recoverability
- ☐ Implementation of Isolation
- ☐ Transaction Definition in SQL
- ☐ Testing for Serializability.

Transaction Concept

- 
- ❑ A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.
 - ❑ E.g., transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
 - ❑ Two main issues to deal with:
 - Failures of various kinds, such as hardware failures and system crashes
 - Concurrent execution of multiple transactions

Required Properties of a Transaction

❑ Consider a transaction to transfer \$50 from account A to account B:

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

❑ **Atomicity requirement**

- If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - Failure could be due to software or hardware
- The system should ensure that updates of a partially executed transaction are not reflected in the database

❑ **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.

Required Properties of a Transaction (Cont.)

☐ **Consistency requirement** in above example:

- The sum of A and B is unchanged by the execution of the transaction

☐ In general, consistency requirements include

- Explicitly specified integrity constraints such as primary keys and foreign keys
- Implicit integrity constraints

e.g., sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand

☐ A transaction, when starting to execute, must see a consistent database.

☐ During transaction execution the database may be temporarily inconsistent.

☐ When the transaction completes successfully the database must be consistent

- Erroneous transaction logic can lead to inconsistency

Required Properties of a Transaction (Cont.)

- ❑ **Isolation requirement** — if between steps 3 and 6 (of the fund transfer transaction), another transaction **T2** is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

T1

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

T2

read(A), read(B), print(A+B)

- ❑ Isolation can be ensured trivially by running transactions **serially**
 - That is, one after the other.
- ❑ However, executing multiple transactions concurrently has significant benefits, as we will see later.

ACID Properties




- ❑ A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:
 - ❑ **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.
 - ❑ **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
 - ❑ **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
 - ❑ **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

Transaction State



- ❑ **Active** – the initial state; the transaction stays in this state while it is executing
- ❑ **Partially committed** – after the final statement has been executed.
- ❑ **Failed** -- after the discovery that normal execution can no longer proceed.
- ❑ **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
 - Restart the transaction
 - can be done only if no internal logical error
 - Kill the transaction
- ❑ **Committed** – after successful completion.

Transaction State (Cont.)



Concurrent Executions



- ❑ Multiple transactions are allowed to run concurrently in the system.

Advantages are:

- **Increased processor and disk utilization**, leading to better transaction *throughput*
 - E.g. one transaction can be using the CPU while another is reading from or writing to the disk
- **Reduced average response time** for transactions: short transactions need not wait behind long ones.

- ❑ **Concurrency control schemes** – mechanisms to achieve isolation

- That is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database
 - Will study in Chapter 15, after studying notion of correctness of concurrent executions.

Schedules



- ❑ **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of all instructions of those transactions
 - Must preserve the order in which the instructions appear in each individual transaction.
- ❑ A transaction that successfully completes its execution will have a **commit** instructions as the last statement
 - By default transaction assumed to execute commit instruction as its last step

Schedule 1



- ❑ Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B .
- ❑ An example of a **serial** schedule in which T_1 is followed by T_2 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

- ❑ A **serial** schedule in which T_2 is followed by T_1 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 3

- Let T_1 and T_2 be the transactions defined previously. The following schedule is not a serial schedule, but it is **equivalent** to Schedule 1.

T_1	T_2
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit

Note -- In schedules 1, 2 and 3, the sum “A + B” is preserved.

- ❑ The following concurrent schedule does not preserve the sum of “ $A + B$ ”

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	 read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Serializability




- ❑ **Basic Assumption** – Each transaction preserves database consistency.
- ❑ Thus, serial execution of a set of transactions preserves database consistency.
- ❑ A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
 1. **conflict serializability**
 2. **view serializability**

Simplified view of transactions



- ☐ We ignore operations other than **read** and **write** instructions
- ☐ We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- ☐ Our simplified schedules consist of only **read** and **write** instructions.

Conflicting Instructions

- 
- ❑ Let l_i and l_j be two Instructions of transactions T_i and T_j respectively.
Instructions l_i and l_j **conflict** if and only if there exists some item Q accessed by both l_i and l_j , and at least one of these instructions wrote Q .
 1. $l_i = \text{read}(Q)$, $l_j = \text{read}(Q)$. l_i and l_j don't conflict.
 2. $l_i = \text{read}(Q)$, $l_j = \text{write}(Q)$. They conflict.
 3. $l_i = \text{write}(Q)$, $l_j = \text{read}(Q)$. They conflict
 4. $l_i = \text{write}(Q)$, $l_j = \text{write}(Q)$. They conflict
 - ❑ Intuitively, a conflict between l_i and l_j forces a (logical) temporal order between them.
 - If l_i and l_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

Conflict Serializability



- ❑ If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent**.
- ❑ We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule

Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6 -- a serial schedule where T_2 follows T_1 , by a series of swaps of non-conflicting instructions. Therefore, Schedule 3 is conflict serializable.

T_1	T_2
read (A) write (A)	
	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 3

T_1	T_2
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6

Conflict Serializability (Cont.)


- ❑ Example of a schedule that is not conflict serializable:

T_3	T_4
read (Q)	write (Q)
write (Q)	

- ❑ We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.




Precedence Graph

- ❑ Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- ❑ **Precedence graph** — a direct graph where the vertices are the transactions (names).
- ❑ We draw an arc from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.
- ❑ We may label the arc by the item that was accessed.
- ❑ **Example**



Testing for Conflict Serializability

- ❑ A schedule is conflict serializable if and only if its precedence graph is acyclic.
- ❑ Cycle-detection algorithms exist which take order n^2 time, where n is the number of vertices in the graph.
 - (Better algorithms take order $n + e$ where e is the number of edges.)
- ❑ If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.
 - That is, a linear order consistent with the partial order of the graph.
 - For example, a serializability order for the schedule (a) would be one of either (b) or (c)



Recoverable Schedules

- ❑ **Recoverable schedule** — if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i **must** appear before the commit operation of T_j .
- ❑ The following schedule is not recoverable if T_9 commits immediately after the read(A) operation.

T_8	T_9
read (A)	
write (A)	
	read (A)
	commit
read (B)	

- ❑ If T_8 should abort, T_9 would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.

Cascading Rollbacks

- ❑ **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

If T_{10} fails, T_{11} and T_{12} must also be rolled back.

- ❑ Can lead to the undoing of a significant amount of work

Cascadeless Schedules

- ❑ **Cascadeless schedules** — for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- ❑ Every cascadeless schedule is also recoverable
- ❑ It is desirable to restrict the schedules to those that are cascadeless
- ❑ Example of a schedule that is NOT cascadeless

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

Concurrency Control



- ❑ A database must provide a mechanism that will ensure that all possible schedules are both:
 - Conflict serializable.
 - Recoverable and preferably cascadeless
- ❑ A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- ❑ Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur
- ❑ Testing a schedule for serializability *after* it has executed is a little too late!
 - Tests for serializability help us understand why a concurrency control protocol is correct
- ❑ **Goal** – to develop concurrency control protocols that will assure serializability.

Weak Levels of Consistency



- ❑ Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - E.g., database statistics computed for query optimization can be approximate (why?)
 - Such transactions need not be serializable with respect to other transactions
- ❑ Tradeoff accuracy for performance

Levels of Consistency in SQL-92



- ☐ **Serializable** — default
- ☐ **Repeatable read** — only committed records to be read, repeated reads of same record must return same value. However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- ☐ **Read committed** — only committed records can be read, but successive reads of record may return different (but committed) values.
- ☐ **Read uncommitted** — even uncommitted records may be read.

- ☐ Lower degrees of consistency useful for gathering approximate information about the database
- ☐ Warning: some database systems do not ensure serializable schedules by default
 - E.g., Oracle and PostgreSQL by default support a level of consistency called snapshot isolation (not part of the SQL standard)

Transaction Definition in SQL



- ❑ Data manipulation language must include a construct for specifying the set of actions that comprise a transaction.
- ❑ In SQL, a transaction begins implicitly.
- ❑ A transaction in SQL ends by:
 - **Commit work** commits current transaction and begins a new one.
 - **Rollback work** causes current transaction to abort.
- ❑ In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
 - Implicit commit can be turned off by a database directive
 - E.g. in JDBC, `connection.setAutoCommit(false);`