

Concurrency Control

Outline



- ☐ Lock-Based Protocols
- ☐ Timestamp-Based Protocols
- ☐ Validation-Based Protocols
- ☐ Multiple Granularity
- ☐ Multiversion Schemes
- ☐ Insert and Delete Operations
- ☐ Concurrency in Index Structures

Lock-Based Protocols



- ❑ A lock is a mechanism to control concurrent access to a data item
- ❑ Data items can be locked in two modes :
 1. *exclusive (X) mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. *shared (S) mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.
- ❑ Lock requests are made to the concurrency-control manager by the programmer. Transaction can proceed only after request is granted.

Lock-Based Protocols (Cont.)

❑ Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- ❑ A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- ❑ Any number of transactions can hold shared locks on an item,
 - But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- ❑ If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.

Lock-Based Protocols (Cont.)

- ❑ Example of a transaction performing locking:

T_2 : **lock-S**(A);
 read (A);
 unlock(A);
 lock-S(B);
 read (B);
 unlock(B);
 display($A+B$)


- ❑ Locking as above is not sufficient to guarantee serializability — if A and B get updated in-between the read of A and B , the displayed sum would be wrong.
- ❑ A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules.

The Two-Phase Locking Protocol



- ❑ This protocol ensures conflict-serializable schedules.
- ❑ Phase 1: Growing Phase
 - Transaction may obtain locks
 - Transaction may not release locks
- ❑ Phase 2: Shrinking Phase
 - Transaction may release locks
 - Transaction may not obtain locks
- ❑ The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).

The Two-Phase Locking Protocol (Cont.)

- 
- ❑ There can be conflict serializable schedules that cannot be obtained if two-phase locking is used.
 - ❑ However, in the absence of extra information (e.g., ordering of access to data), two-phase locking is needed for conflict serializability in the following sense:
 - Given a transaction T_i that does not follow two-phase locking, we can find a transaction T_j that uses two-phase locking, and a schedule for T_i and T_j that is not conflict serializable.

Lock Conversions



- ❑ Two-phase locking with lock conversions:
 - First Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- ❑ This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.

Automatic Acquisition of Locks

- ❑ A transaction T_i issues the standard read/write instruction, without explicit locking calls.
- ❑ The operation **read**(D) is processed as:
 - if** T_i has a lock on D
 - then**
 - read(D)
 - else begin**
 - if necessary wait until no other transaction has a **lock-X** on D
 - grant T_i a **lock-S** on D ;
 - read(D)
 - end**

Automatic Acquisition of Locks (Cont.)

❑ **write(D)** is processed as:

if T_i has a **lock-X** on D

then

 write(D)

else begin

 if necessary wait until no other transaction has any lock on D ,

if T_i has a **lock-S** on D

then

upgrade lock on D to **lock-X**

else

 grant T_i a **lock-X** on D

 write(D)

end;

❑ All locks are released after commit or abort

Deadlocks

- ❑ Consider the partial schedule

T_3	T_4
lock-x (B)	
read (B)	
$B := B - 50$	
write (B)	
	lock-s (A)
	read (A)
	lock-s (B)
lock-x (A)	

- ❑ Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- ❑ Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.

Deadlocks (Cont.)



- ❑ Two-phase locking *does not* ensure freedom from deadlocks.
- ❑ In addition to deadlocks, there is a possibility of **starvation**.
- ❑ **Starvation** occurs if the concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- ❑ Concurrency control manager can be designed to prevent starvation.

Deadlocks (Cont.)




- ❑ The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- ❑ When a deadlock occurs there is a possibility of cascading roll-backs.
- ❑ Cascading roll-back is possible under two-phase locking. To avoid this, follow a modified protocol called **strict two-phase locking** -- a transaction must hold all its exclusive locks till it commits/aborts.
- ❑ **Rigorous two-phase locking** is even stricter. Here, *all* locks are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.

Implementation of Locking



- ❑ A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- ❑ The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- ❑ The requesting transaction waits until its request is answered
- ❑ The lock manager maintains a data-structure called a **lock table** to record granted locks and pending requests
- ❑ The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked

Lock Table




- ❑ Dark blue rectangles indicate granted locks; light blue indicate waiting requests
- ❑ Lock table also records the type of lock granted or requested
- ❑ New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- ❑ Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- ❑ If transaction aborts, all waiting or granted requests of the transaction are deleted
 - lock manager may keep a list of locks held by each transaction, to implement this efficiently

Deadlock Handling




- ❑ System is deadlocked if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set.
- ❑ **Deadlock prevention** protocols ensure that the system will *never* enter into a deadlock state. Some prevention strategies :
 - Require that each transaction locks all its data items before it begins execution (predeclaration).
 - Impose partial ordering of all data items and require that a transaction can lock data items only in the order specified by the partial order.

More Deadlock Prevention Strategies

- 
- ❑ Following schemes use transaction timestamps for the sake of deadlock prevention alone.
 - ❑ **wait-die** scheme — non-preemptive
 - older transaction may wait for younger one to release data item. (older means smaller timestamp) Younger transactions never wait for older ones; they are rolled back instead.
 - a transaction may die several times before acquiring needed data item
 - ❑ **wound-wait** scheme — preemptive
 - older transaction *wounds* (forces rollback) of younger transaction instead of waiting for it. Younger transactions may wait for older ones.
 - may be fewer rollbacks than *wait-die* scheme.


Deadlock prevention (Cont.)

- 
- ❑ Both in *wait-die* and in *wound-wait* schemes, a rolled back transactions is restarted with its original timestamp. Older transactions thus have precedence over newer ones, and starvation is hence avoided.
 - ❑ **Timeout-Based Schemes:**
 - a transaction waits for a lock only for a specified amount of time. If the lock has not been granted within that time, the transaction is rolled back and restarted,
 - Thus, deadlocks are not possible
 - simple to implement; but starvation is possible. Also difficult to determine good value of the timeout interval.


Deadlock Detection

- ❑ Deadlocks can be described as a *wait-for graph*, which consists of a pair $G = (V, E)$,
 - V is a set of vertices (all the transactions in the system)
 - E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$.
- ❑ If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j , implying that T_i is waiting for T_j to release a data item.
- ❑ When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted in the wait-for graph. This edge is removed only when T_j is no longer holding a data item needed by T_i .
- ❑ The system is in a deadlock state if and only if the wait-for graph has a cycle. Must invoke a deadlock-detection algorithm periodically to look for cycles.

Deadlock Detection (Cont.)



□ Wait-for graph without a cycle



□ Wait-for graph with a cycle

Deadlock Recovery

❑ When deadlock is detected :

- Some transaction will have to rolled back (made a victim) to break deadlock. Select that transaction as victim that will incur minimum cost.
- Rollback -- determine how far to roll back transaction
 - **Total rollback:** Abort the transaction and then restart it.
 - More effective to roll back transaction only as far as necessary to break deadlock.
- Starvation happens if same transaction is always chosen as victim. Include the number of rollbacks in the cost factor to avoid starvation